

Anonymous^a

ABSTRACT

This supplementary document provides the base-model checkpoint identifiers (Table S1) for the three backbone models used in this study. All three backbones were loaded without an explicit `revision=` argument; the SHA-1 hashes below are the snapshots present in the HuggingFace cache at the time of each experiment and are sufficient to retrieve the identical weights.

Table S1

HuggingFace Hub repository identifiers and cache-snapshot SHA-1 hashes for the three base-model checkpoints used in this study. Weights can be retrieved via `from_pretrained(repo_id, revision=sha1)`. [†]Loaded without `revision=`; SHA-1 inferred from cache state at execution time (see provenance note above).

Backbone	SHA-1 (40 chars)	Repository
Qwen 2.5-1.5B-Instruct	989aa798...aa306	Qwen/Qwen2.5-1.5B-Instruct
Gemma 3 1B IT	dcc83ea8...78752	google/gemma-3-1b-it
Gemma 4 E2B IT [†]	70af34e2...22a03	google/gemma-4-E2B-it

Table S1: Base-Model Checkpoint Identifiers

All three backbones were loaded from the Hugging Face Hub using `AutoModelForCausalLM.from_pretrained(repo_id)` without an explicit `revision=` argument. The SHA-1 hashes below identify the snapshots present in the local cache at the time of each experiment and are sufficient to retrieve the identical weights via `from_pretrained(repo_id, revision=sha1)`.

Provenance note for Gemma 4 E2B IT. Snapshot 70af34e was downloaded on 2026-06-26; the Gemma 4 experiments ran on 2026-06-27, when this was the only snapshot in cache. A second snapshot (3e22461) entered the cache on 2026-09-30 and was not used in any reported experiment.

Full 40-character SHA-1 strings:

- Qwen 2.5-1.5B-Instruct: 989aa7980e4cf806f80c7fef2b1adb7bc71aa306
- Gemma 3 1B IT: dcc83ea841ab6100d6b47a070329e1ba4cf78752
- Gemma 4 E2B IT: 70af34e20bd4b7a91f0de6b22675850c43922a03

These identifiers are also listed in `configs/checkpoints.yaml` of the replication package (Zenodo DOI: [10.5281/zenodo.23146342](https://doi.org/10.5281/zenodo.23146342)).